

Import the database file directly into a Pandas dataframe object. Execute the code cell exactly as written.

```
In [38]: import pandas as pd
url = "http://54.243.252.9/engr-1330-webroot/4-Databases/EcommerceCustomers.csv"
dataset = pd.read_csv(url)
```

Problem 1 (5 pts) : *Import any remaining libraries and describe the dataframe

```
In [39]: #GIVE YOUR ANSWER FOR Problem-1 IN THIS CELL
import matplotlib.pyplot as plt
import numpy as np
import math
import seaborn as sns

dataset.describe()
```

```
Out[39]:
```

	Avg. Session Length	Time on App	Time on Website	Length of Membership	Yearly Amount Spent
count	500.000000	500.000000	500.00000	500.000000	500.000000
mean	33.052400	12.052600	37.06100	3.532800	499.320000
std	0.992767	0.995588	1.01109	0.999581	79.318045
min	29.500000	8.500000	33.90000	0.300000	257.000000
25%	32.300000	11.400000	36.37500	2.900000	445.000000
50%	33.100000	12.000000	37.10000	3.500000	499.000000
75%	33.700000	12.800000	37.70000	4.100000	549.250000
max	36.100000	15.100000	40.00000	6.900000	766.000000

The dataframe lists the email, Address, and avatar of

ecommerce customers and lists each persons information

namely the average time they play, time on the app,

time on their website, how long theyve been a member,

and the yearly amount of money they spend on their app

with this information tons of data can be analyzed

to maximize profits.

Problem 2: *We will be working with an ecommerce dataset to predict the customer's yearly spending amount in dollars based on 4 features, namely, average session length in minutes (average in-store style advice sessions), time on app in minutes, time on website in minutes, length of membership in years. Specific tasks for this exercise are given below.*

Task-1 (5 pts): Read the `EcommerceCustomers.csv` file and store it in a variable named `ecommerce`. Display the first 5 rows of the dataframe `ecommerce`.

****Problem 0.1 above should automatically load the file into the object `dataset`. If that fails, the CSV file is located at http://54.243.252.9/engr-1330-webroot/4-Databases/EcommerceCustomers.csv****

In [40]: *#GIVE YOUR ANSWER FOR TASK-1 IN THIS CELL*

```
ecommerce = pd.read_csv('EcommerceCustomers.csv')
ecommerce.head()
```

Out[40]:

	Email	Address	Avatar	Avg. Session Length	Time on App	Time on Website	Length of Membership
0	mstephenson@fernandez.com	835 Frank Tunnel\nWrightmouth, MI 82180-9605	Violet	34.5	12.7	39.6	
1	hduke@hotmail.com	4547 Archer Common\nDiazchester, CA 06566-8576	DarkGreen	31.9	11.1	37.3	
2	pallen@yahoo.com	24645 Valerie Unions Suite 582\nCobbborough, D...	Bisque	33.0	11.3	37.1	

	Email	Address	Avatar	Avg. Session Length	Time on App	Time on Website	Length of Membership
3	riverarebecca@gmail.com	1414 David Throughway\nPort Jason, OH 22070-1220	SaddleBrown	34.3	13.7	36.7	
4	mstephens@davidson-herman.com	14023 Rodriguez Passage\nPort Jacobville, PR 3...	MediumAquaMarine	33.3	12.8	37.5	

Task-2 (5 pts): Get the information (i.e, the number of rows and the data types present in each column) and the basic statistical measures about the dataframe **ecommerce** using the appropriate functions.

```
In [41]: #GIVE YOUR ANSWER FOR TASK-2 IN THIS CELL
print('The total number of rows are',ecommerce.shape[0], 'And the total number of columns are',ecommerce.shape[1])

dataType = ecommerce.dtypes
print('=====')
print('These are the data types of each column')
print(dataType)
```

The total number of rows are 500 And the total number of columns are 8 !

=====


These are the data types of each column

```
Email          object
Address         object
Avatar          object
Avg. Session Length  float64
Time on App      float64
Time on Website  float64
Length of Membership float64
Yearly Amount Spent int64
dtype: object
```

Task-3 (5 pts): Display the plot between all parameters in the dataset using the appropriate plotting function in the Seaborn library.


```
In [42]: #GIVE YOUR ANSWER FOR TASK-3 IN THIS CELL

sns.pairplot(ecommerce)
plt.show()
```



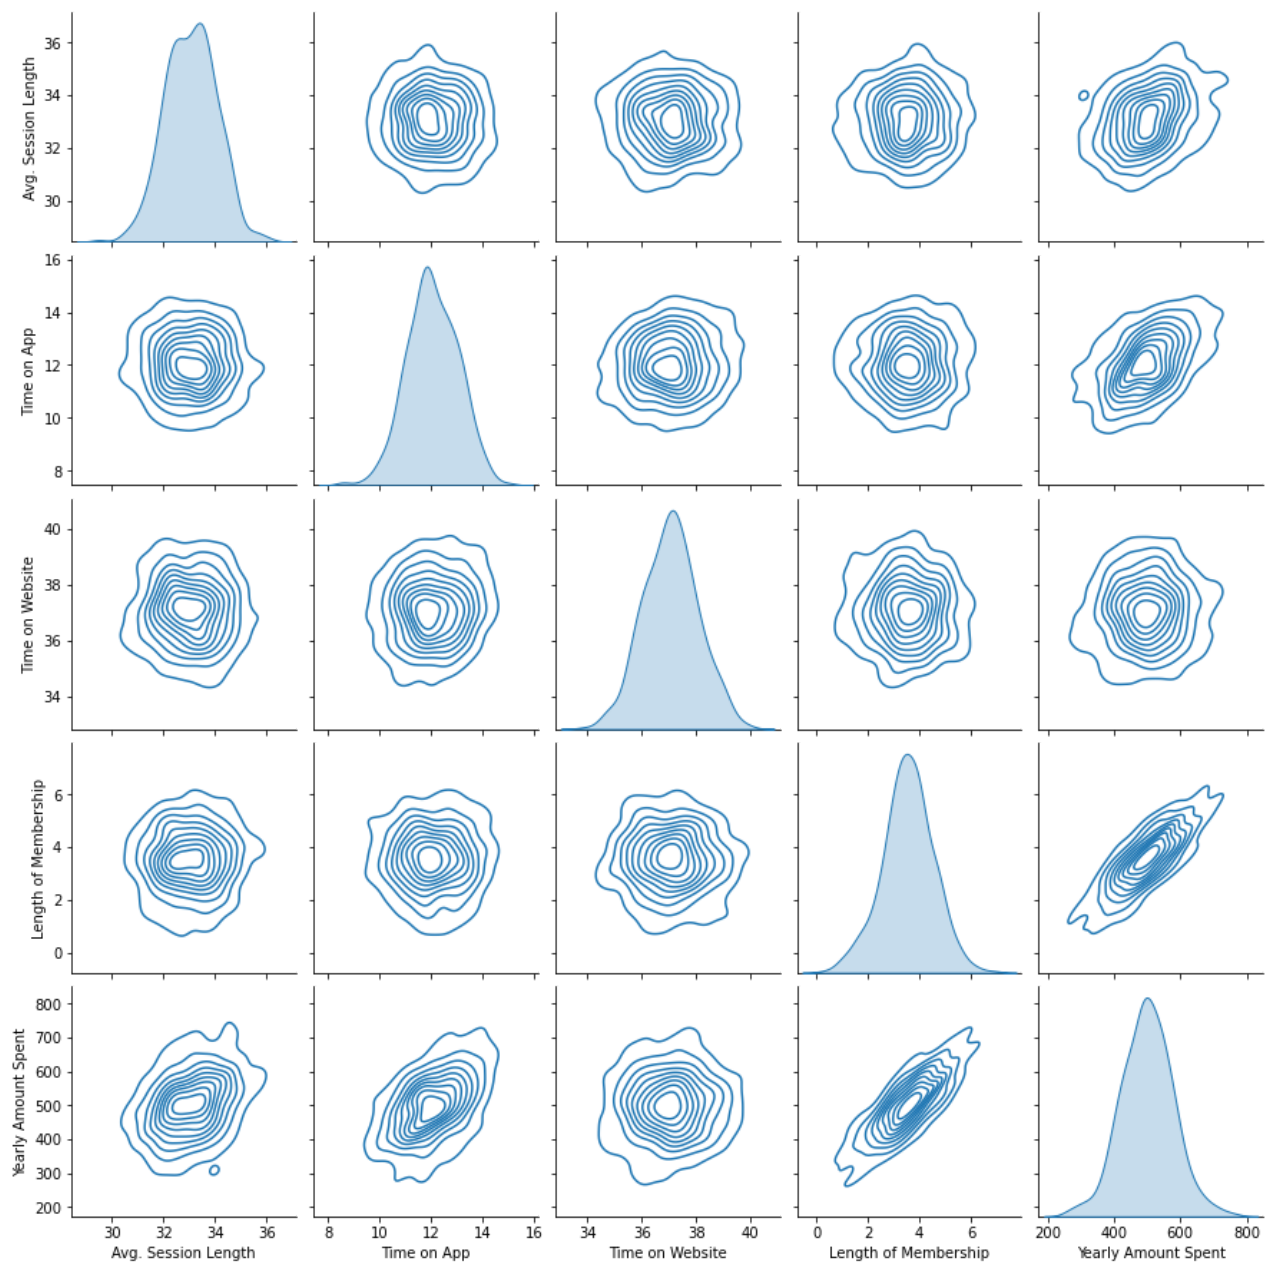
```
In [9]: sns.pairplot(ecommerce, diag_kind="hist")
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x1367d842d60>
```



```
In [10]: sns.pairplot(ecommerce, kind="kde")
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x1367e1595e0>
```

I added several models for completion

Task-4: (5 pts) Compute the correlation coefficient between all parameters in the dataset using the appropriate function in the Pandas library and store it in a variable named `corr`. Display the heat map of correlation values stored in `corr` using the appropriate function in the Seaborn library. The heat map must display an appropriate title. Identify if there is any colinearity between the input parameters.


```
In [11]: #GIVE YOUR ANSWER FOR TASK-4 IN THIS CELL
#make this more general, find the correlation matrix. generate the heatmap
corr = ecommerce.corr()
print(corr)

sns.heatmap(ecommerce.corr());
```

	Avg. Session Length	Time on App	Time on Website	\
Avg. Session Length	1.000000	-0.026253	-0.034575	

Time on App	-0.026253	1.000000	0.084760
Time on Website	-0.034575	0.084760	1.000000
Length of Membership	0.061958	0.028690	-0.049037
Yearly Amount Spent	0.356523	0.497717	-0.003118

	Length of Membership	Yearly Amount Spent
Avg. Session Length	0.061958	0.356523
Time on App	0.028690	0.497717
Time on Website	-0.049037	-0.003118
Length of Membership	1.000000	0.809364
Yearly Amount Spent	0.809364	1.000000



Task-5: (5 pts) Index and pick the 4 feature columns and store it in a variable named `x` . Similarly, index and pick the column named **Yearly Amount Spent** and store it in a variable named `y` . Print the first 5 rows of `x` and `y` to the output.

```
In [20]: #GIVE YOUR ANSWER FOR TASK-5 IN THIS CELL
y = ecommerce['Yearly Amount Spent']
x = ecommerce.drop(columns=['Yearly Amount Spent', 'Email', 'Address', 'Avatar'], axis=1)
#x = ecommerce.drop(columns=['Email'], axis=1)
#x = ecommerce.drop(columns=['Address'], axis=1)
#x = ecommerce.drop(columns=['Avatar'], axis=1)

print(x)
print('=====')
print(y)
```

	Avg. Session Length	Time on App	Time on Website	Length of Membership
0	34.5	12.7	39.6	4.1
1	31.9	11.1	37.3	2.7
2	33.0	11.3	37.1	4.1
3	34.3	13.7	36.7	3.1
4	33.3	12.8	37.5	4.4
...
495	33.2	13.6	36.4	3.7
496	34.7	11.7	37.2	3.6
497	32.6	11.5	38.3	5.0

498	33.3	12.4	36.8	2.3
499	33.7	12.4	35.8	2.7

[500 rows x 4 columns]

=====

```
0      588
1      392
2      488
3      582
4      599
```

...

```
495     574
496     529
497     552
498     456
499     498
```

Name: Yearly Amount Spent, Length: 500, dtype: int64

Task-6 (5 pts): Randomly (without replacement) pick 66.66% of the rows from the dataframe **data** and store it in a variable named **training_set** . Similarly, randomly (without replacement) pick 33.33% of the rows from the dataframe **data** and store it in a variable named **test_set** .

```
In [31]: #GIVE YOUR ANSWER FOR TASK-6 IN THIS CELL
from sklearn import svm
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.linear_model import LinearRegression

x_train = x.sample(frac=.6666, replace = False)
y_train = y.sample(frac=.6666, replace = False)
x_test = x.sample(frac=.3333, replace = False)
y_test = y.sample(frac=.3333, replace = False)
```

Task-7 (5 pts): Train the model using **x_train** and **y_train** . Print the coefficients and the intercept to the output.

```
In [32]: # Task 7 goes here

k = LinearRegression()
print(k.fit(x_train,y_train))
print('-----')
print(k.coef_)
print(k.intercept_)

LinearRegression()
-----
[-2.2189509  4.19521231  3.18499642  0.73577379]
406.0245688833435
```

```
In [34]: # More in depth analysis

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.6666, random_stat

k = LinearRegression()
k.fit(x_train,y_train)
k.coef_
```

```

predic = k.predict(x_test)

sns.scatterplot(y_test, predic)

mae = metrics.mean_absolute_error(y_test, predic)
mse = metrics.mean_squared_error(y_test, predic)
rmse = np.sqrt(metrics.mean_squared_error(y_test, predic))
r2 = metrics.r2_score(y_test, predic)
lin = LinearRegression()


print('Mean Absolute error is:', mae )
print('Mean Square error is:', mse)
print('Root Mean square error is:', rmse)
print('R2 Score is:', r2)

```

Mean Absolute error is: 7.891975911150439
Mean Square error is: 99.803890664589
Root Mean square error is: 9.990189721150895
R2 Score is: 0.9850960607668636

C:\Users\medra\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



Now for .33 version

```

In [35]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3333, random_stat

k = LinearRegression()
k.fit(x_train,y_train)
k.coef_

predic = k.predict(x_test)


```

```
sns.scatterplot(y_test, predic)
```

C:\Users\medra\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[35]: <AxesSubplot:xlabel='Yearly Amount Spent'>



```
In [36]: mae = metrics.mean_absolute_error(y_test, predic)
mse = metrics.mean_squared_error(y_test, predic)
rmse = np.sqrt(metrics.mean_squared_error(y_test, predic))
r2 = metrics.r2_score(y_test, predic)

print('Mean Absolute error is:', mae )
print('Mean Square error is:', mse)
print('Root Mean square error is:', rmse)
print('R2 Score is:', r2)
```

```
Mean Absolute error is: 7.474829475033762
Mean Square error is: 86.34983305468205
Root Mean square error is: 9.292461087068487
R2 Score is: 0.9874765294165274
```

In []: